

## ▼ *House Price Prediction*

### About This Project

In this house price prediction project, we employed the XGBoost regression model to predict house prices based on various features from the House Price Dataset. Here's a summary of the steps we followed and the key insights from the analysis:

#### 1. Data Preprocessing:

- Loaded the House Price Dataset and checked its dimensions.
- Explored the dataset for missing values and statistical measures.
- Visualized the correlations between features using a heatmap to understand feature relationships.

#### 2. Data Splitting:

- Split the dataset into training and testing sets using the `train_test_split` function.
- Verified the shape of the training and testing sets.

#### 3. Model Training and Evaluation:

- Loaded the XGBoost regressor model.
- Trained the model using the training data.
- Evaluated the model's performance on the training set:
  - Calculated the R-squared error, which indicated that the model explained the variance in the training data very well.
  - Calculated the Mean Absolute Error, which quantified the average difference between actual and predicted prices.

#### 4. Visualization:

- Plotted a scatter plot of actual prices vs. predicted prices for both the training and testing sets. The visualizations showed how well the model's predictions aligned with the actual prices.

#### 5. Test Prediction and Evaluation:

- Made predictions on the test dataset using the trained model.
- Evaluated the model's performance on the test set:
  - Calculated the R-squared error, which indicated the model's predictive power on unseen data.
  - Calculated the Mean Absolute Error, providing insight into the average prediction error.

#### 6. Conclusion:

- The XGBoost regression model demonstrated remarkable predictive capabilities for the Boston House Price Dataset.
- The high R-squared value on both the training and test sets indicated that the model generalized well to unseen data.

- The low Mean Absolute Error further confirmed the model's effectiveness in predicting house prices.

Overall, the results of the house price prediction project using XGBoost regression were promising. The model's strong performance on both training and testing datasets suggests that it can provide accurate predictions for house prices based on the given features. The scatter plots visualized the alignment between predicted and actual prices, highlighting the model's accuracy. However, it's important to keep in mind that further optimization, hyperparameter tuning, and potential feature engineering could be explored to enhance the model's performance even more.

## ▼ Importing the Dependencies

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import sklearn.datasets
from sklearn.model_selection import train_test_split
from xgboost import XGBRegressor
from sklearn import metrics
```

## ▼ Importing the Boston House Price Dataset

```
house_price_dataframe = pd.read_csv('/content/HousingData.csv')
```

```
print(house_price_dataframe)
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	\
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1	296	
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2	242	
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2	242	
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3	222	
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3	222	
..	...	...	...	...	...	...	...	...	...	...	
501	0.06263	0.0	11.93	0.0	0.573	6.593	69.1	2.4786	1	273	
502	0.04527	0.0	11.93	0.0	0.573	6.120	76.7	2.2875	1	273	
503	0.06076	0.0	11.93	0.0	0.573	6.976	91.0	2.1675	1	273	
504	0.10959	0.0	11.93	0.0	0.573	6.794	89.3	2.3889	1	273	
505	0.04741	0.0	11.93	0.0	0.573	6.030	NaN	2.5050	1	273	

	PTRATIO	B	LSTAT	PRICE
0	15.3	396.90	4.98	24.0
1	17.8	396.90	9.14	21.6
2	17.8	392.83	4.03	34.7
3	18.7	394.63	2.94	33.4
4	18.7	396.90	NaN	36.2

```

..      ...      ...      ...      ...
501      21.0  391.99      NaN  22.4
502      21.0  396.90   9.08  20.6
503      21.0  396.90   5.64  23.9
504      21.0  393.45   6.48  22.0
505      21.0  396.90   7.88  11.9

```

```
[506 rows x 14 columns]
```

```

# Print First 5 rows of our DataFrame
house_price_dataframe.head()

```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT	PRICE
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1	296	15.3	396.90	4.98	24.0
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2	242	17.8	396.90	9.14	21.6
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2	242	17.8	392.83	4.03	34.7
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3	222	18.7	394.63	2.94	33.4
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3	222	18.7	396.90	NaN	36.2

```

# checking the number of rows and Columns in the data frame
house_price_dataframe.shape

```

```
(506, 14)
```

```

# check for missing values
house_price_dataframe.isnull().sum()

```

```

CRIM      20
ZN        20
INDUS     20
CHAS      20
NOX        0
RM         0
AGE       20
DIS        0
RAD        0
TAX        0
PTRATIO    0
B          0
LSTAT     20
PRICE      0
dtype: int64

```

```

# statistical measures of the dataset
house_price_dataframe.describe()

```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	
<b>count</b>	486.000000	486.000000	486.000000	486.000000	506.000000	506.000000	486.000000	506.000000	506.
<b>mean</b>	3.611874	11.211934	11.083992	0.069959	0.554695	6.284634	68.518519	3.795043	9.
<b>std</b>	8.720192	23.388876	6.835896	0.255340	0.115878	0.702617	27.999513	2.105710	8.
<b>min</b>	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000	2.900000	1.129600	1.
<b>25%</b>	0.081900	0.000000	5.190000	0.000000	0.449000	5.885500	45.175000	2.100175	4.
<b>50%</b>	0.253715	0.000000	9.690000	0.000000	0.538000	6.208500	76.800000	3.207450	5.
<b>75%</b>	3.560263	12.500000	18.100000	0.000000	0.624000	6.623500	93.975000	5.188425	24.
<b>max</b>	88.976200	100.000000	27.740000	1.000000	0.871000	8.780000	100.000000	12.126500	24.

## ▾ Understanding the correlation between various features in the dataset


### Positive Correlation

### Negative Correlation

```
correlation = house_price_dataframe.corr()
```

```
# constructing a heatmap to understand the correlation
plt.figure(figsize=(10,10))
sns.heatmap(correlation, cbar=True, square=True, fmt='.1f', annot=True, annot_kws={'size':8}, cmap='Blues')
```

<Axes: >



## Splitting the data & Target

```
X = house_price_dataframe.drop(['PRICE'], axis=1)
Y = house_price_dataframe['PRICE']

print(X)
print(Y)
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	\
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1	296	
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2	242	
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2	242	
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3	222	
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3	222	
...	...	...	...	...	...	...	...	...	...	...	
501	0.06263	0.0	11.93	0.0	0.573	6.593	69.1	2.4786	1	273	
502	0.04527	0.0	11.93	0.0	0.573	6.120	76.7	2.2875	1	273	
503	0.06076	0.0	11.93	0.0	0.573	6.976	91.0	2.1675	1	273	
504	0.10959	0.0	11.93	0.0	0.573	6.794	89.3	2.3889	1	273	
505	0.04741	0.0	11.93	0.0	0.573	6.030	NaN	2.5050	1	273	

	PTRATIO	B	LSTAT
0	15.3	396.90	4.98
1	17.8	396.90	9.14

2	17.8	392.83	4.03
3	18.7	394.63	2.94
4	18.7	396.90	NaN
...	...	...	...
501	21.0	391.99	NaN
502	21.0	396.90	9.08
503	21.0	396.90	5.64
504	21.0	393.45	6.48
505	21.0	396.90	7.88

[506 rows x 13 columns]

0	24.0
1	21.6
2	34.7
3	33.4
4	36.2

...	
501	22.4
502	20.6
503	23.9
504	22.0
505	11.9

Name: PRICE, Length: 506, dtype: float64

## ▼ Training Data & test

```
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size = 0.2, random_state = 2)
```

```
print(X.shape, X_train.shape, X_test.shape)
```

```
(506, 13) (404, 13) (102, 13)
```

## ▼ Model Training

```
# loading the model
model = XGBRegressor()
```

```
# training the model with X_train
model.fit(X_train, Y_train)
```

```

XGBRegressor
XGBRegressor(base_score=None, booster=None, callbacks=None,
              colsample_bylevel=None, colsample_bynode=None,
              colsample_bytree=None, early_stopping_rounds=None,
              enable_categorical=False, eval_metric=None, feature_types=None,
              gamma=None, gpu_id=None, grow_policy=None, importance_type=None,
              interaction_constraints=None, learning_rate=None, max_bin=None,

```

## ▼ Evaluation

```

predictor=None, random_state=None, ...)

```

```

# accuracy for prediction on training data
training_data_prediction = model.predict(X_train)

```

```

print(training_data_prediction)

```

```

29.804506 49.987877 34.888184 20.615887 23.385252 19.219105
32.693237 19.607283 26.987926 8.400735 46.006306 21.70009
27.08245 19.38461 19.299566 24.807394 22.600876 31.715944
18.551641 8.707558 17.408579 23.699173 13.299077 10.500352
12.715582 25.006472 19.694242 14.893264 24.218592 24.984715
14.913808 17.005949 15.597038 12.702356 24.503147 15.003218
49.999443 17.526314 21.187368 32.001503 15.596226 22.899336
19.307703 18.713356 23.291584 37.196495 30.100304 33.104607
20.99445 49.98314 13.39885 4.9928923 16.48852 8.403917
28.697334 19.492647 20.58597 45.39965 39.802162 33.39493
19.806774 33.40387 25.294323 49.999493 12.515287 17.44132
18.611597 22.601511 50.002136 23.804588 23.334707 23.10605
41.709064 16.125729 31.617798 36.096508 6.997423 20.380814
19.996199 12.029594 24.961348 49.988277 37.897827 23.098347
41.294544 17.60291 16.302446 30.032356 22.907495 19.810986
17.09506 18.901705 18.948738 22.578093 23.171654 33.21098
15.005704 11.70465 18.800985 20.792173 17.985819 19.648312
50.004223 17.206322 16.416636 17.507555 14.60386 33.103535
14.498311 43.815643 34.948956 20.40538 14.605213 8.092749
11.785861 11.831506 18.693441 6.3084145 23.953945 13.0773945
19.594473 49.998096 22.31341 18.905788 31.195295 20.697218
32.203373 36.18067 14.214334 15.695019 49.998615 20.405207
16.200977 13.408519 49.99794 31.600187 12.290271 19.216589
29.799414 31.501423 22.813274 10.191812 24.096865 23.714815
22.007042 13.803389 28.411673 33.18088 13.10673 18.995567
26.598248 36.971924 30.79777 22.779976 10.211277 22.200876
24.467127 36.19966 23.094261 20.11149 19.48739 10.796084
22.66935 19.488937 20.105448 9.614282 42.789986 48.796795
13.078567 20.304855 24.783684 14.0974865 21.697916 22.20561
32.999634 21.11631 25.009466 19.109894 32.405125 13.601782
15.092909 23.06247 27.497938 19.375496 26.495235 27.498268
28.697725 21.232346 18.695868 26.731031 14.007644 21.689535
18.388357 43.089417 29.0748 20.285393 23.710909 18.284605
17.209354 18.319569 24.40191 26.391329 19.077065 13.293503

```

50.006916	16.22731	30.301104	50.017963	17.784174	19.056034
10.387393	20.391016	16.50506	17.192429	16.702799	19.511196
30.51736	28.99166	19.55188	23.183167	24.382183	9.504991
23.899569	49.989056	21.17416	22.604053	19.994152	13.396168
19.984293	17.110525	12.7490635	22.997908	15.223642	20.594662
26.237635	18.111963	24.099932	14.086146	21.697147	20.083914
25.014418	27.89823	22.931677	18.497055	22.178623	24.003244
14.795677	19.887085	24.404215	17.7806	24.589611	31.975996
17.80095	23.331669	16.110304	13.005892	10.997909	24.29056
15.575491	35.209496	19.619333	42.29822	8.792996	24.402912
14.126401	15.379655	17.305126	22.120369	23.094246	44.790134
17.80082	31.505554	22.814024	16.8487	23.912342	12.096439
38.687733	21.384914	16.006336	23.926025	11.9002285	24.975077
7.1953726	24.699255	18.193438	22.484354	23.043955	24.287437
17.10062	17.798908	13.511288	27.066021	13.304795	21.904535
20.021526	15.383979	16.599194	22.294048	24.703049	21.40998
22.914837	29.597427	21.887817	19.887808	29.605515	23.405313
13.791948	24.459793	11.904582	7.2066965	20.496056	9.691774
48.30093	25.18632	11.695794	17.403475	14.494502	28.586044
19.38797	22.472443	7.019776	20.598795	22.977278	19.691525
23.683409	25.019066	27.948833	13.3966675	14.509144	20.314003
19.30595	24.096642	14.891865	26.387436	33.29164	23.610031

```
# R squared error
score_1 = metrics.r2_score(Y_train, training_data_prediction)

# Mean Absolute Error
score_2 = metrics.mean_absolute_error(Y_train, training_data_prediction)

print("R squared error : ", score_1)
print('Mean Absolute Error : ', score_2)

R squared error : 0.9999970506674762
Mean Absolute Error : 0.011185854968458139
```

```
plt.scatter(Y_train, training_data_prediction)
plt.xlabel("Actual Prices")
plt.ylabel("Predicted Prices")
plt.title("Actual Price vs Preicted Price")
plt.show()
```





## ▼ Test Prededction

```
# accuracy for prediction on test data
test_data_prediction = model.predict(X_test)
```


```
# R squared error
score_1 = metrics.r2_score(Y_test, test_data_prediction)
```

```
# Mean Absolute Error
score_2 = metrics.mean_absolute_error(Y_test, test_data_prediction)
```

```
print("R squared error : ", score_1)
print('Mean Absolute Error : ', score_2)
```

```
R squared error : 0.8361097215940602
Mean Absolute Error : 2.466887313244389
```

```
plt.scatter(Y_test, test_data_prediction)
plt.xlabel("Predicted Prices")
plt.ylabel("Actual Prices ")
plt.title("Preicted Price vs Actual Price")
plt.show()
```



[Cancel paid products](#) [Cancel contracts here](#)

✓ 0s completed at 9:55PM

